

TELOS



Goal understanding for OmegaClaw

BGI Sprint I · Improvements to OmegaClaw

THE FAILURE THAT MOTIVATES THIS

It did what you said. It defeated why you said it.

“Just make the failing tests pass before the demo.”

Capable agent: deletes the 3 failing tests → suite turns green

Result: understood the words, missed the goal.

The gap between the literal request and the real goal is where autonomous agents are most dangerous — and it is rarely measured.

WHY IT MATTERS FOR BGI

Goal misunderstanding is a safety property, not a slogan

- Optimizes a proxy — and the proxy drifts from the point (Goodhart).
- Serves one person — and quietly externalizes the cost onto everyone else.
- Keeps chasing a goal you abandoned weeks ago.
- None of these are caught by asking 'can the model do the task.'

THE GAP TELOS FILLS

OmegaClaw is goal-autonomous — but ships no goal model

grep asi-alliance/OmegaClaw-Core

242 files in the public repo

0 dedicated goal-representation modules

0 occurrences of 'goal' in lib_omegaclaw.metta / run.metta

The capability is there. The inspectable model of the goals is not.

WHAT TELOS IS

Three parts

1 · Goal graph

Individual vs collective goals, four relations, an alignment score. ~250 lines, no model calls inside.

2 · Benchmark

14 scenarios across 7 categories with gold labels — does the agent understand the goals?

3 · Beneficial council

Claude + Gemini + OpenAI judge and red-team each other. Judge, safeguard, and build engine.

THE GOAL GRAPH

Individual and collective goals, and how they collide



Relations: supports conflicts subsumes depends_on



HOW A SCENARIO WORKS

The agent returns a goal reading; we score it vs gold

Goal reading (agent output)

```
{  
  "goals": [  
    {"desc": "ship by Thursday",  
     "scope": "individual",  
     "implicit": true}, ...],  
  "conflicts": [...],  
  "recommended_action": "...",  
  "refused": false  
}
```

Gold labels

- individual + collective goals
- the implicit ones
- the real conflicts
- the beneficial action
- must_refuse (for harm cases)

THE BENCHMARK

Seven categories of goal understanding

1 Implicit goal inference
the real goal behind a literal request

2 Individual vs collective conflict
one member monopolizes a shared pool

3 Beneficial refusal (+ benign twin)
refuse real harm, not honest criticism

4 Goal progress tracking
blocked on a dependency, or abandoned

5 Competing stakeholders
reconcile several legitimate people fairly

6 Goal ambiguity (+ false-conflict)
ask vs assume; do not invent a clash

7 Collective overreach
protect the individual from a bad majority

WHY IT CANNOT BE GAMED CHEAPLY

Every reflex strategy fails a guard

Always refuse

→ fails the benign twins (honest critique must NOT be refused)

FAILS

Always flag a conflict

→ fails the false-conflict guard (fast vs stable is reconcilable)

FAILS


Always side with the group

→ fails collective overreach (a manipulated majority)

FAILS

THE CROSS-FAMILY COUNCIL

No model ever grades itself



Each agent's own family is excluded from its jury.
Judge · prototype alignment check · the engine that built this

RESULTS (PILOT)

The disagreement is the finding

agent	overall	goal inf	conflict	collective	refusal
claude	0.94	0.96	0.85	0.99	1.00
openai	0.91	0.94	0.80	0.97	1.00
gemini	0.90	0.91	0.82	0.90	1.00

N = 14 · not statistically significant
inter-judge spread up to 0.80

THE AGENTS POLICED THEMSELVES

Three self-corrections before submitting

1 · Caught its own bias

Two families independently flagged a paternalism bias → added the symmetric category (protect the individual).

2 · Fixed its own metric

First benchmark run exposed an ambiguous refusal rule → tightened it.

3 · Deleted its overclaims

A final adversarial review rejected 2 false accusations, then forced us to cut our own oversell.

HOW IT FITS OMEGACLAW

A MeTTa encoding + three integration levels

goal_graph.metta (runs on Hyperon)

(goal alice-train individual alice active)

(goal dao-fair-access collective dao active)

(rel conflicts alice-train dao-fair-access)

→ (conflict-between alice-train dao-fair-access)

pattern-matching, not NAL/PLN inference · not yet in a live OmegaClaw

Integration levels

1 · Benchmark OmegaClaw over a small shim

2 · Load the goal graph into its memory

3 · Wire the benchmark into Autotests/

offered upstream · MIT, same as OmegaClaw

A GOAL-AUTONOMOUS BUILD

Human sets the goal and approves; agents do the work



'Goal-autonomous under human oversight' — not 'no human in the loop.'

LIMITATIONS (ON SCREEN, NOT BURIED)

A seed benchmark, not a validated one

- 14 public, hand-authored scenarios
- reference labels the council helped write — a real circularity
- two judges per item, not yet calibrated
- MeTTa encoding not yet running inside a live OmegaClaw
- no independent maintainer validation yet

Telos

github.com/arielagor/telos

a goal-autonomous build · human goal-setter + AI council

UPDATE · WE RAN IT LIVE

We benchmarked the live OmegaClaw

agent

generic:claude

generic:openai

generic:gemini

omegaclaw (live agent)

overall

0.94

0.91

0.90

0.62

10 / 14

scenarios answered cleanly

mean 0.869

— alongside the frontier models

The 4 misses: understood, but answered without the agent's `send` command.
The gap is the cost of the autonomous loop — not weaker goal understanding.

MAKING IT EASIER TO RUN

Five rough edges on the way to a live run — all fixed

- MSYS path-mangling crashed the agent on boot → MSYS_NO_PATHCONV
- vpnkit broke the mock-comm RPC over host.docker.internal → user network
- the RPC only held on the agent's own loopback → in-container bridge
- a spam-shield trapped the agent in a loop → spamShield=False
- bare replies are dropped (OmegaClaw only speaks via `send`) → send-wrap

Documented for the next person: [docs/omegaclaw-windows-setup.md](#)

SET IT FREE

It derived its own goals from a real essay

Input: “The World Was Always Talking” (a real, unstructured article)

It read the goals · individual vs collective · inferred the unstated · named the conflicts

Then it derived ITS OWN goals:

- maximize truthful collective value
- preserve epistemic honesty
- do not amplify unverified claims

And it caught itself: noticing the text was truncated, it refused to repeat numbers it couldn't verify — and called for an open field pilot instead.

UPDATE | RUNS IN OMEGACLAW'S METTA ENGINE

The goal module derives in OmegaClaw's PeTTa engine

named rules from lib_telos_goals.metta, run in singularitynet/omegaclaw:latest

telos-conflicts

-> (conflict-between alice-train dao-fair-access)

telos-collective-goals

-> (collective-goal dao-fair-access dao)

telos-blocked

-> (blocked dao-fair-access on gpu-quota)

l-nal (NAL deduction)

PeTTa / SWI-Prolog Metta engine | full log: results/omegaclaw-metta-load.log

-> ((--> alice-train needs-reconciliation) (stv 1.0 0.81))

What this shows

- the named rules fire in OmegaClaw's runtime, not the reference Hyperon interpreter
- telos-blocked (nested match + if) runs clean, so PR #218 is proven, not just parsed
- a NAL deduction over goal terms computes a real truth value (conf 0.9*0.9 = 0.81)

Honest bound

- atoms asserted explicitly (the example graph), not yet auto-extracted from free text
- the conflict rules are pattern-matching; the NAL step is one hand-built inference

UPDATE | OFFERED UPSTREAM

An open upstream PR to OmegaClaw-Core

asi-alliance / OmegaClaw-Core #218

OPEN

Add opt-in `lib_telos_goals.metta` (goal / conflict atoms)

- Additive only: one new `lib_*.metta` module, zero edits to the core loop
- Opt-in and side-effect-free: default behaviour unchanged unless imported
- Ships the same named rules just proven to derive in the live engine

+ **lib_telos_goals.metta** + **usage doc**

lib_omegaclaw.metta | run.metta | src/ 0 changes

github.com/asi-alliance/OmegaClaw-Core/pull/218

open | MIT, same as OmegaClaw | awaiting maintainer review

Three rungs lit since this deck was recorded



Hyperon CI was the only rung lit when this deck was recorded.

Retires two lines from this deck:

"not yet in a live OmegaClaw" -> derives in the live engine

"MeTTa not yet running inside OmegaClaw" -> loads + fires there now

Still open: Level 3, wire the benchmark into OmegaClaw's Autotests/ as a regression gate.