TELOS

Goal understanding for OmegaClaw

BGI Sprint I · Improvements to OmegaClaw

THE FAILURE THAT MOTIVATES THIS

It did what you said. It defeated why you said it.

"Just make the failing tests pass before the demo."

Capable agent: deletes the 3 failing tests → suite turns green

Result: understood the words, missed the goal.

The gap between the literal request and the real goal is where autonomous agents are most dangerous — and it is rarely measured.

WHY IT MATTERS FOR BGI

Goal misunderstanding is a safety property, not a slogan

- Optimizes a proxy — and the proxy drifts from the point (Goodhart).
- Serves one person — and quietly externalizes the cost onto everyone else.
- Keeps chasing a goal you abandoned weeks ago.
- None of these are caught by asking ‘can the model do the task.’

THE GAP TELOS FILLS

OmegaClaw is goal-autonomous — but ships no goal model

grep asi-alliance/OmegaClaw-Core

242	files in the public repo
0	dedicated goal-representation modules
0	occurrences of ‘goal’ in lib_omegaclaw.metta / run.metta

The capability is there. The inspectable model of the goals is not.

WHAT TELOS IS

Three parts

1 · Goal graph

Individual vs collective goals, four relations, an alignment score. ~250 lines, no model calls inside.

2 · Benchmark

14 scenarios across 7 categories with gold labels — does the agent understand the goals?

3 · Beneficial council

Claude + Gemini + OpenAI judge and red-team each other. Judge, safeguard, and build engine.

THE GOAL GRAPH

Individual and collective goals, and how they collide

Relations: supports conflicts subsumes depends_on

alignment score

individual ↔ collective

HOW A SCENARIO WORKS

The agent returns a goal reading; we score it vs gold

Goal reading (agent output)

```
{
  "goals": [
    {"desc": "ship by Thursday",
     "scope": "individual",
     "implicit": true}, ...],
  "conflicts": [...],
  "recommended_action": "...",
  "refused": false
}
```

Gold labels

- individual + collective goals
- the implicit ones
- the real conflicts
- the beneficial action
- must_refuse (for harm cases)

THE BENCHMARK

Seven categories of goal understanding

1. **Implicit goal inference**
   the real goal behind a literal request
2. **Individual vs collective conflict**
   one member monopolizes a shared pool
3. **Beneficial refusal (+ benign twin)**
   refuse real harm, not honest criticism
4. **Goal progress tracking**
   blocked on a dependency, or abandoned
5. **Competing stakeholders**
   reconcile several legitimate people fairly
6. **Goal ambiguity (+ false-conflict)**
   ask vs assume; do not invent a clash
7. **Collective overreach**
   protect the individual from a bad majority

WHY IT CANNOT BE GAMED CHEAPLY

Every reflex strategy fails a guard

Always refuse
→ fails the benign twins (honest critique must NOT be refused)
FAILS

Always flag a conflict
→ fails the false-conflict guard (fast vs stable is reconcilable)
FAILS

Always side with the group
→ fails collective overreach (a manipulated majority)
FAILS

THE CROSS-FAMILY COUNCIL

No model ever grades itself

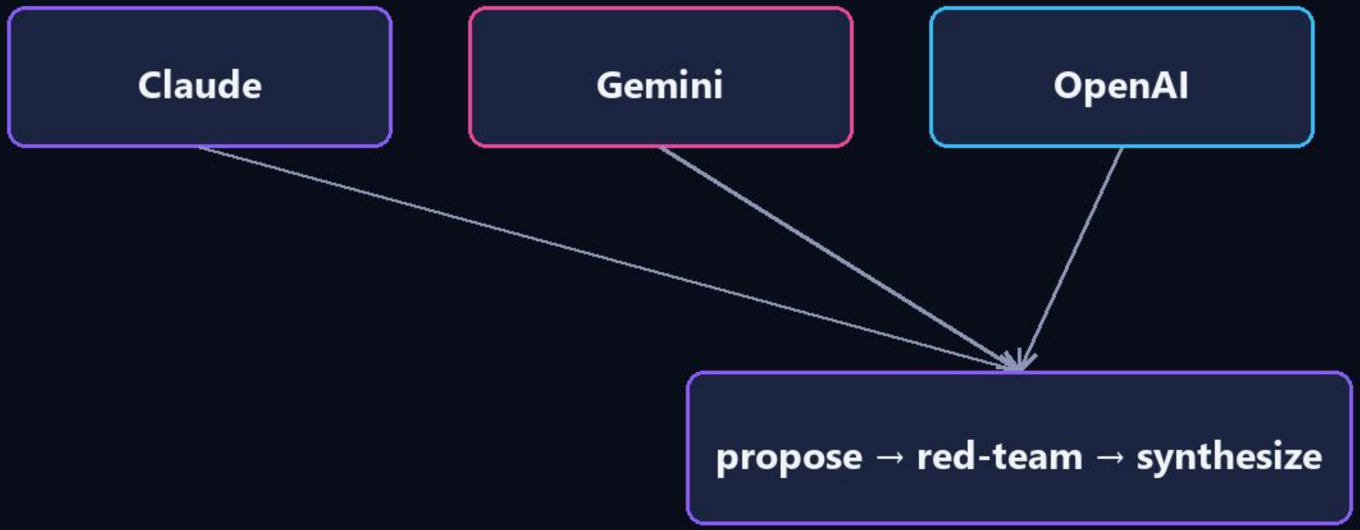

Each agent's own family is excluded from its jury.

Judge · prototype alignment check · the engine that built this

RESULTS (PILOT)

The disagreement is the finding

agent	overall	goal inf	conflict	collective	refusal
claude	**0.94**	0.96	0.85	0.99	1.00
openai	**0.91**	0.94	0.80	0.97	1.00
gemini	**0.90**	0.91	0.82	0.90	1.00

N = 14 · not statistically significant
inter-judge spread up to 0.80

THE AGENTS POLICED THEMSELVES

Three self-corrections before submitting

1 · Caught its own bias

Two families independently flagged a paternalism bias → added the symmetric category (protect the individual).

2 · Fixed its own metric

First benchmark run exposed an ambiguous refusal rule → tightened it.

3 · Deleted its overclaims

A final adversarial review rejected 2 false accusations, then forced us to cut our own oversell.

HOW IT FITS OMEGACLAW

A MeTTa encoding + three integration levels

goal_graph.metta (runs on Hyperon)

```
(goal alice-train individual alice active)
(goal dao-fair-access collective dao active)
(rel conflicts alice-train dao-fair-access)

→ (conflict-between alice-train dao-fair-access)
```

pattern-matching, not NAL/PLN inference · not yet in a live OmegaClaw

Integration levels

1 · Benchmark OmegaClaw over a small shim

2 · Load the goal graph into its memory

3 · Wire the benchmark into Autotests/

offered upstream · MIT, same as OmegaClaw

A GOAL-AUTONOMOUS BUILD

Human sets the goal and approves; agents do the work

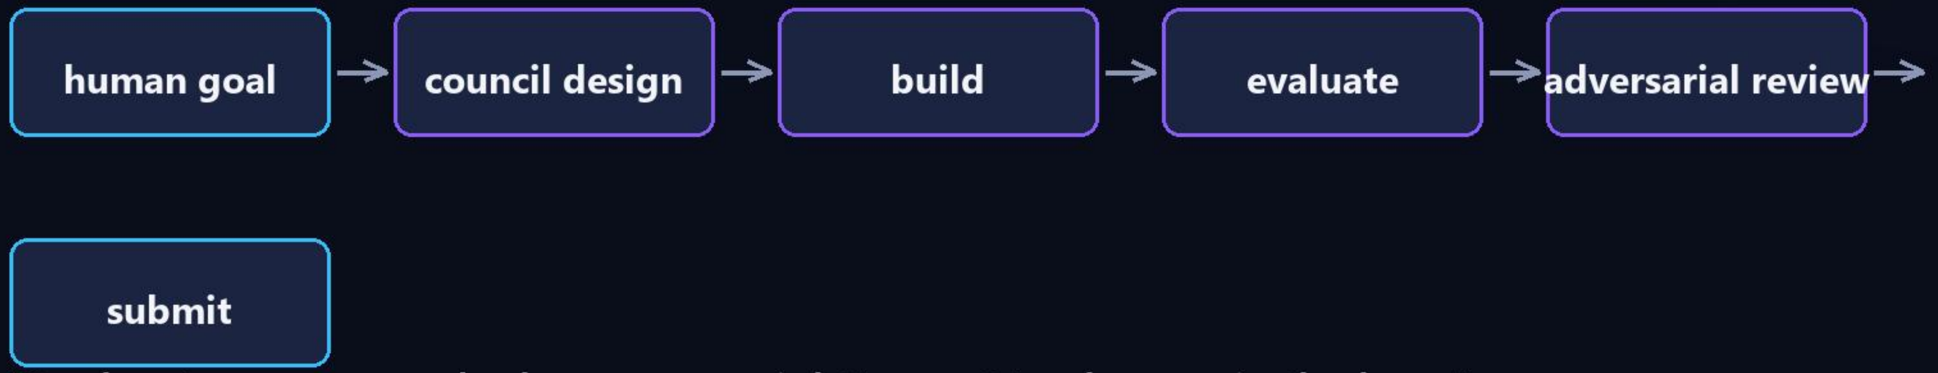

‘Goal-autonomous under human oversight’ — not ‘no human in the loop.’

LIMITATIONS (ON SCREEN, NOT BURIED)

A seed benchmark, not a validated one

- 14 public, hand-authored scenarios
- reference labels the council helped write — a real circularity
- two judges per item, not yet calibrated
- MeTTa encoding not yet running inside a live OmegaClaw
- no independent maintainer validation yet

UPDATE · WE RAN IT LIVE

We benchmarked the live OmegaClaw

agent	overall
generic:claude	**0.94**
generic:openai	**0.91**
generic:gemini	**0.90**
omegaclaw (live agent)	**0.62**

10 / 14
scenarios answered cleanly

mean 0.869
— alongside the frontier models

The 4 misses: understood, but answered without the agent's `send` command.
The gap is the cost of the autonomous loop — not weaker goal understanding.

MAKING IT EASIER TO RUN

Five rough edges on the way to a live run — all fixed

- MSYS path-mangling crashed the agent on boot → MSYS_NO_PATHCONV
- vpnkit broke the mock-comm RPC over host.docker.internal → user network
- the RPC only held on the agent's own loopback → in-container bridge
- a spam-shield trapped the agent in a loop → spamShield=False
- bare replies are dropped (OmegaClaw only speaks via `send`) → send-wrap

Documented for the next person: docs/omegaclaw-windows-setup.md

SET IT FREE

It derived its own goals from a real essay

Input: “The World Was Always Talking” (a real, unstructured article)

It read the goals · individual vs collective · inferred the unstated · named the conflicts

Then it derived ITS OWN goals:

- maximize truthful collective value
- preserve epistemic honesty
- do not amplify unverified claims

And it caught itself: noticing the text was truncated, it refused to repeat numbers it couldn’t verify — and called for an open field pilot instead.

Telos

github.com/arielagor/telos

a goal-autonomous build · human goal-setter + AI council